

# Udiddit, a social news aggregator

## Introduction

Udiddit, a social news aggregation, web content rating, and discussion website, is currently using a risky and unreliable Postgres database schema to store the forum posts, discussions, and votes made by their users about different topics.

The schema allows posts to be created by registered users on certain topics, and can include a URL or a text content. It also allows registered users to cast an upvote (like) or downvote (dislike) for any forum post that has been created. In addition to this, the schema also allows registered users to add comments on posts.

Here is the DDL used to create the schema:

```
CREATE TABLE bad_posts (  
    id SERIAL PRIMARY KEY,  
    topic VARCHAR(50),  
    username VARCHAR(50),  
    title VARCHAR(150),  
    url VARCHAR(4000) DEFAULT NULL,  
    text_content TEXT DEFAULT NULL,  
    upvotes TEXT,  
    downvotes TEXT  
);  
  
CREATE TABLE bad_comments (  
    id SERIAL PRIMARY KEY,  
    username VARCHAR(50),  
    post_id BIGINT,  
    text_content TEXT  
);
```

## Part I: Investigate the existing schema

As a first step, investigate this schema and some of the sample data in the project's SQL workspace. Then, in your own words, outline three (3) specific things that could be improved about this schema. Don't hesitate to outline more if you want to stand out!

1. In bad\_posts table :
  - upvotes and downvotes should be INTEGER, not TEXT
  - Add NOT NULL constraints on username, topic and title
  - username and topic should be UNIQUE
2. In bad\_comments table :
  - Add foreign key constraint to link post\_id to posts
  - Add NOT NULL constraints on username, post\_id and text\_content
  - Usernames should be UNIQUE
3. Add indexes on commonly queried fields
4. Add timestamp fields for created\_at/updated\_at for posts and comments
5. Add more tables e.g. for users, topics to appropriately structure the data

## Part II: Create the DDL for your new schema

Having done this initial investigation and assessment, your next goal is to dive deep into the heart of the problem and create a new schema for Udiddit. Your new schema should at least reflect fixes to the shortcomings you pointed to in the previous exercise. To help you create the new schema, a few guidelines are provided to you:

1. Guideline #1: here is a list of features and specifications that Udiddit needs in order to support its website and administrative interface:
  - a. Allow new users to register:
    - i. Each username has to be unique
    - ii. Usernames can be composed of at most 25 characters
    - iii. Usernames can't be empty
    - iv. We won't worry about user passwords for this project
  - b. Allow registered users to create new topics:
    - i. Topic names have to be unique.
    - ii. The topic's name is at most 30 characters
    - iii. The topic's name can't be empty
    - iv. Topics can have an optional description of at most 500 characters.
  - c. Allow registered users to create new posts on existing topics:
    - i. Posts have a required title of at most 100 characters
    - ii. The title of a post can't be empty.
    - iii. Posts should contain either a URL or a text content, **but not both**.
    - iv. If a topic gets deleted, all the posts associated with it should be automatically deleted too.
    - v. If the user who created the post gets deleted, then the post will remain, but it will become dissociated from that user.
  - d. Allow registered users to comment on existing posts:
    - i. A comment's text content can't be empty.
    - ii. Contrary to the current linear comments, the new structure should allow comment threads at arbitrary levels.
    - iii. If a post gets deleted, all comments associated with it should be automatically deleted too.
    - iv. If the user who created the comment gets deleted, then the comment will remain, but it will become dissociated from that user.
    - v. If a comment gets deleted, then all its descendants in the thread structure should be automatically deleted too.
  - e. Make sure that a given user can only vote once on a given post:
    - i. Hint: you can store the (up/down) value of the vote as the values 1 and -1 respectively.
    - ii. If the user who cast a vote gets deleted, then all their votes will remain, but will become dissociated from the user.

- iii. If a post gets deleted, then all the votes for that post should be automatically deleted too.
2. Guideline #2: here is a list of queries that Udiddit needs in order to support its website and administrative interface. Note that you don't need to produce the DQL for those queries: they are only provided to guide the design of your new database schema.
- a. List all users who haven't logged in in the last year.
  - b. List all users who haven't created any post.
  - c. Find a user by their username.
  - d. List all topics that don't have any posts.
  - e. Find a topic by its name.
  - f. List the latest 20 posts for a given topic.
  - g. List the latest 20 posts made by a given user.
  - h. Find all posts that link to a specific URL, for moderation purposes.
  - i. List all the top-level comments (those that don't have a parent comment) for a given post.
  - j. List all the direct children of a parent comment.
  - k. List the latest 20 comments made by a given user.
  - l. Compute the score of a post, defined as the difference between the number of upvotes and the number of downvotes
3. Guideline #3: you'll need to use normalization, various constraints, as well as indexes in your new database schema. You should use named constraints and indexes to make your schema cleaner.
4. Guideline #4: your new database schema will be composed of five (5) tables that should have an auto-incrementing id as their primary key.

Once you've taken the time to think about your new schema, write the DDL for it in the space provided here:

```
-- create users table
CREATE TABLE users (
  id SERIAL PRIMARY KEY,
  username VARCHAR(25) NOT NULL UNIQUE,
  last_login TIMESTAMP DEFAULT CURRENT_TIMESTAMP,

  CONSTRAINT user_name_not_empty CHECK (LENGTH(TRIM(username)) > 0)
);
```

```

--create topics table
CREATE TABLE topics(
    id SERIAL PRIMARY KEY,
    topic_name VARCHAR(30) NOT NULL UNIQUE,
    topic_description VARCHAR(500),

    CONSTRAINT topic_name_not_empty CHECK (LENGTH(TRIM(topic_name)) > 0)
);

-- create posts table
CREATE TABLE posts(
    id SERIAL PRIMARY KEY,
    title VARCHAR(100) NOT NULL,
    topic_id INTEGER NOT NULL ,
    user_id INTEGER ,
    url VARCHAR(2000),
    text_content TEXT,
    created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,

    CONSTRAINT post_title_not_empty CHECK (LENGTH(TRIM(title)) > 0),
    CONSTRAINT either_ur_or_text_content CHECK ( (url IS NOT NULL AND
text_content IS NULL) OR (url IS NULL AND text_content IS NOT NULL) ),
    CONSTRAINT fk_topic FOREIGN KEY (topic_id) REFERENCES topics(id) ON
DELETE CASCADE, -- topic_deleted_post_deleted
    CONSTRAINT fk_user FOREIGN KEY (user_id) REFERENCES users(id) ON
DELETE SET NULL -- user_deleted_post_remain
);

-- create comments table
CREATE TABLE comments(
    id SERIAL PRIMARY KEY,
    post_id INTEGER NOT NULL,
    user_id INTEGER ,
    parent_comment_id INTEGER ,
    text_content TEXT NOT NULL,

```

```

        created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,

        CONSTRAINT comment_not_empty CHECK (LENGTH(TRIM(text_content)) > 0),
        CONSTRAINT fk_post FOREIGN KEY (post_id) REFERENCES posts(id) ON
DELETE CASCADE, -- post_deleted_comments_deleted
        CONSTRAINT fk_user FOREIGN KEY (user_id) REFERENCES users(id) ON
DELETE SET NULL, -- user_deleted_comment_remain
        CONSTRAINT fk_parent_comment FOREIGN KEY (parent_comment_id)
REFERENCES comments(id) ON DELETE CASCADE --
comment_deleted_decendants_deleted
);

--create votes table
CREATE TABLE votes(
    id SERIAL PRIMARY KEY,
    post_id INTEGER NOT NULL,
    user_id INTEGER ,
    vote_value INTEGER NOT NULL,

    CONSTRAINT vote_value_check CHECK (vote_value IN (-1,1)), -- 1
upvote, -1 downvote
    CONSTRAINT fk_post FOREIGN KEY (post_id) REFERENCES posts(id) ON
DELETE CASCADE, -- post_deleted_votes_deleted
    CONSTRAINT fk_user FOREIGN KEY (user_id) REFERENCES users(id) ON
DELETE SET NULL, -- user_deleted_votes_remain
    CONSTRAINT one_vote_per_user_per_post UNIQUE (post_id, user_id)
);

-- create indexes
CREATE INDEX idx_users_last_login ON users(last_login);

CREATE INDEX idx_votes_posts ON votes(post_id);
CREATE INDEX idx_votes_users ON votes(user_id);

CREATE INDEX idx_comments_post_id ON comments(post_id);
CREATE INDEX idx_comments_users_id ON comments(user_id);

```

```
CREATE INDEX idx_comments_parent_id ON comments (parent_comment_id);  
CREATE INDEX idx_comments_created_at ON comments(created_at);  
  
CREATE INDEX idx_posts_user_id ON posts(user_id);  
CREATE INDEX idx_posts_topic_id ON posts(topic_id);  
CREATE INDEX idx_posts_url ON posts(url);  
CREATE INDEX idx_posts_created_at ON posts(created_at);
```

## Part III: Migrate the provided data

Now that your new schema is created, it's time to migrate the data from the provided schema in the project's SQL Workspace to your own schema. This will allow you to review some DML and DQL concepts, as you'll be using INSERT...SELECT queries to do so. Here are a few guidelines to help you in this process:

1. Topic descriptions can all be empty
2. Since the bad\_comments table doesn't have the threading feature, you can migrate all comments as top-level comments, i.e. without a parent
3. You can use the Postgres string function **regexp\_split\_to\_table** to unwind the comma-separated votes values into separate rows
4. Don't forget that some users only vote or comment, and haven't created any posts. You'll have to create those users too.
5. The order of your migrations matter! For example, since posts depend on users and topics, you'll have to migrate the latter first.
6. Tip: You can start by running only SELECTs to fine-tune your queries, and use a LIMIT to avoid large data sets. Once you know you have the correct query, you can then run your full INSERT...SELECT query.
7. **NOTE:** The data in your SQL Workspace contains thousands of posts and comments. The DML queries may take at least 10-15 seconds to run.

Write the DML to migrate the current data in bad\_posts and bad\_comments to your new database schema:

```
-- migrate unique users from all sources
INSERT INTO users(username) (
    SELECT DISTINCT username
    FROM (
        SELECT username FROM bad_posts
        UNION
        SELECT username FROM bad_comments
        UNION
        SELECT DISTINCT regexp_split_to_table(upvotes,',') AS username
    FROM bad_posts
    UNION
        SELECT DISTINCT regexp_split_to_table(downvotes,',') AS username
    FROM bad_posts
```



```

    ) all_usernames
    WHERE username IS NOT NULL
);

-- migrate unique topics from bad_posts
INSERT INTO topics(topic_name) (
    SELECT DISTINCT topic
    FROM bad_posts
    WHERE topic IS NOT NULL
);

INSERT INTO posts(title, topic_id, user_id, url, text_content) (
    SELECT
        SUBSTRING(bp.title, 1, 100) as title,
        t.id AS topic_id,
        u.id AS user_id,
        bp.url,
        bp.text_content
    FROM bad_posts bp
    JOIN topics t ON t.topic_name = bp.topic
    LEFT JOIN users u ON u.username = bp.username
);

-- migrate comments
INSERT INTO comments(post_id, user_id, text_content) (
    SELECT
        bc.post_id,
        u.id AS user_id,
        bc.text_content
    FROM bad_comments bc
    LEFT JOIN users u ON u.username = bc.username
);

-- migrate upvotes
INSERT INTO votes(post_id, user_id, vote_value) (

```

```

SELECT
    bp.id AS post_id,
    u.id AS user_id,
    1 AS vote_value
FROM bad_posts bp
    CROSS JOIN regexp_split_to_table(bp.upvotes, ',') AS
upvoter(username)
    JOIN users u ON u.username = upvoter.username
);
-- migrate downvotes
INSERT INTO votes(post_id, user_id, vote_value) (
    SELECT
        bp.id AS post_id,
        u.id AS user_id,
        -1 AS vote_value
    FROM bad_posts bp
        CROSS JOIN regexp_split_to_table(bp.downvotes, ',') AS
downvoter(username)
        JOIN users u ON u.username = downvoter.username
);

DROP TABLE bad_comments;
DROP TABLE bad_posts;

```